

CockroachDB

Distributed SQL at scale



THE TWELVE FACTORS

I. Codebase

One codebase tracked in revision control, many deploys

II. Dependencies

Explicitly declare and isolate dependencies

III. Config

Store config in the environment

IV. Backing services

Treat backing services as attached resources

V. Build, release, run

Strictly separate build and run stages

VI. Processes

Execute the app as one or more stateless processes

VII. Port binding

Export services via port binding

VIII. Concurrency

Scale out via the process model

IX. Disposability

Maximize robustness with fast startup and graceful shutdown

X. Dev/prod parity

Keep development, staging, and production as similar as possible

XI. Logs

Treat logs as event streams

XII. Admin processes

Run admin/management tasks as one-off processes

The underlying
philosophy for
microservices is a
12 factor app.

12factor.net



Factor Number VI (Processes):
Make your
microservices
stateless.



Our micro-
services should
be **Cattle**



not **Pets**



But...

Where should we
keep our data?



Traditional
databases are
single instance:
if it goes down,
YOU go down.



And if you want to scale it, you have to build that single instance taller and taller.



Not horizontally,
like our
microservices.



"I'VE LOOKED AT THE SYSTEM AND FOUND A
SINGLE POINT OF FAILURE. YOU."

And...
single-instance is a
SPOF.



To get resiliency
and scalability, you
can use a **NoSQL**
Database...

But can it be **SQL**?





Distributed SQL blends old and new:

- NoSQL scalability
- SQL compatible
- ACID guarantees

A C I D
ATOMICITY CONSISTENCY ISOLATION DURABILITY



The usual suspects
have cloud
offerings.



But this also needs to be solved across clouds or in our own labs and datacenters to avoid cloud lock-in.



kubernetes

Most
importantly, it
should run
inside K8s.



CockroachDB
turns SQL into a
resilient,
scalable,
ACID-compliant
cluster.

Why name it CockroachDB?



CockroachDB in a sentence...



Cockroach DB

CockroachDB clusterizes a single-instance SQL database so it can survive failures, scales, yet act as if it's a single, logical database.

What CockroachDB Solves



Examples of what CockroachDB can solve...

1. Provide a resilient datastore for microservices which runs in Kubernetes.
2. Host transactional, always-on database workloads.
3. Provide systems of record (ledger, transactional backends, & more.) or systems of access (such as identity access systems.)
4. Run multi-region databases with immediate consistency (full ACID guarantees).
5. Meet GDPR requirements by choosing where to domicile data.

Production Examples



CockroachDB in Production

1. **Telecom provider:** Customer service virtual agent and chat.
2. **Major Bank:** Core services in a multi-region, resilient cluster.
3. **European badging system manufacturer:** Access management system requiring low-latency reads and immediate consistency.

Production Examples



CockroachDB in Production

4. **Financial Company:** Customer Identity Access Management with immediate consistency for a company that provides many online products with a single login. Prior replicated system had production-impacting time lags.
5. **Gaming Company:** Primary transactional database for financial transactions across all of Europe while adhering to the GDPR regulations.

CockroachDB Characteristics




Characteristics

1. **OLTP Database** best for transactional workloads.
2. **ANSI SQL Postgres** wire-compatible (9.6).
3. **ACID Compliant** with Serializable Isolation.
4. **Postgres Driver** connect to any node to connect and your code sees a single DB.
5. **Horizontally scalable** by adding nodes whenever capacity or resiliency is needed.
6. **Kubernetes friendly** by just installing a cluster with a helm chart. To scale, just add instances.

CockroachDB: a unique distributed architecture

self contained, aware nodes participate in global cluster

Each node within a cluster is self-contained and has locational awareness of their self and others



- Every node is a CONSISTENT gateway to the entire database
- Intelligence packed with data
 - Management & optimization
 - Standard SQL engine
 - Enterprise security
 - Ecosystem integration

CockroachDB: a unique distributed architecture

global database cluster coordination and logic

Spin up a node anywhere (public and private clouds)
and then point it at the cluster, which takes care of:

- Coordination & consensus for queries/transactions
- Replication, repair & rebalancing of data across cluster upon addition/removal of nodes
- Attach location to any data to set domiciling & replication constraints

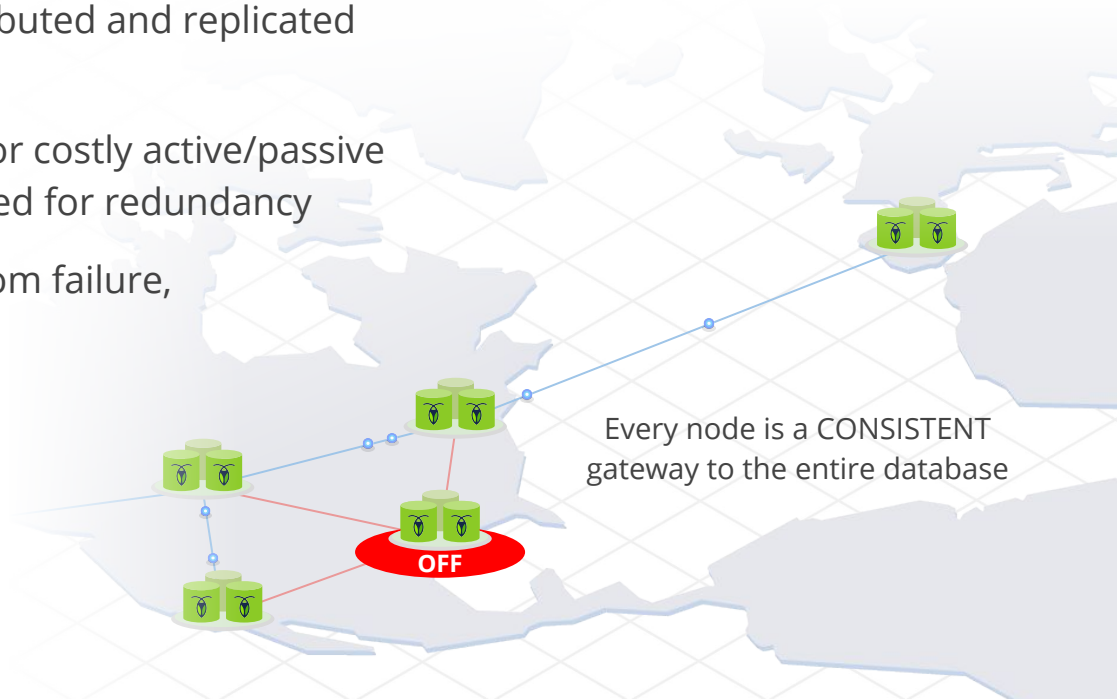
Inherently **multi-cloud**



CockroachDB: Always on and naturally **resilient**

Your data is always on and always available

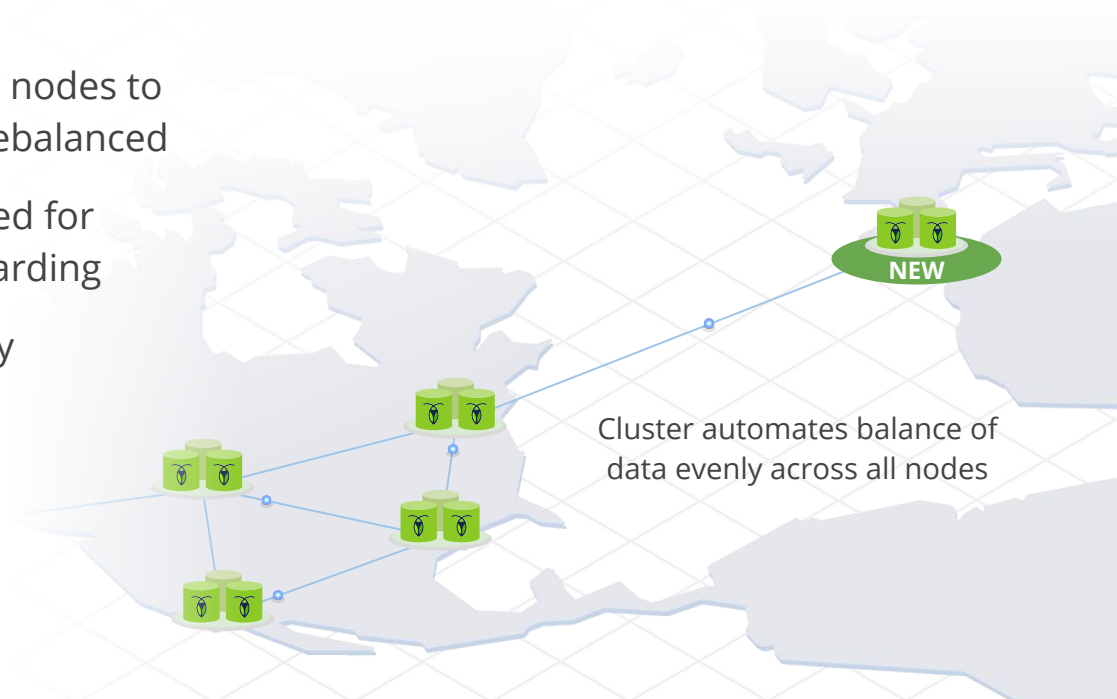
1. On failure, data is efficiently redistributed and replicated across nodes within clusters
2. CockroachDB eliminates the need for costly active/passive or complex CDC architectures needed for redundancy
3. Minimize impact & recovery time from failure, with Cockroach RTO is zero



CockroachDB: **Scale** your data not your complexity

Replication, repair & rebalancing of data across cluster upon addition or removal of nodes

1. To expand capacity, simply add new nodes to the cluster & data is automatically rebalanced
2. Automated balancing eliminates need for manual sharding and complex resharding
3. Balancing optimizes server efficiency (storage and compute)




CockroachDB: Guarantees global consistency

Ensures consistency across distributed transactions

CockroachDB uses clocks and concurrency controls to deliver **full ACID** transactions at scale even in a distributed environment

Serializable isolation protects from write skew and dirty reads

CockroachDB: Serializable isolation
in a distributed SQL database



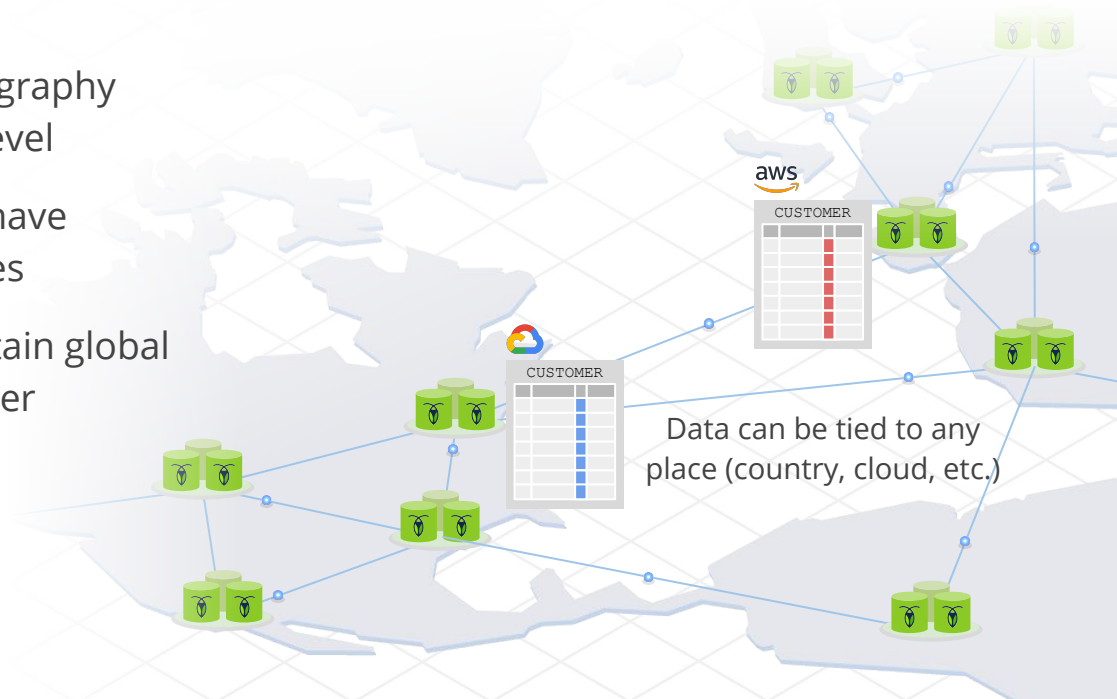
*transactions may not physically execute serialized in time, rather they execute as if they have. They are guaranteed to appear serialized



CockroachDB: Tie your data to a **location**

Geo-partition your data to set
domiciling & replication constraints

1. Tie explicit “ranges” of data to a geography or any address at the table or row level
2. Comply with privacy regulation OR have data follow a user to reduce latencies
3. Tie data to explicit clouds and maintain global access to all nodes throughout cluster



CockroachDB: Inherently **multi-cloud**

Implement a globally consistent database
across clouds and even on premise



aws



Azure



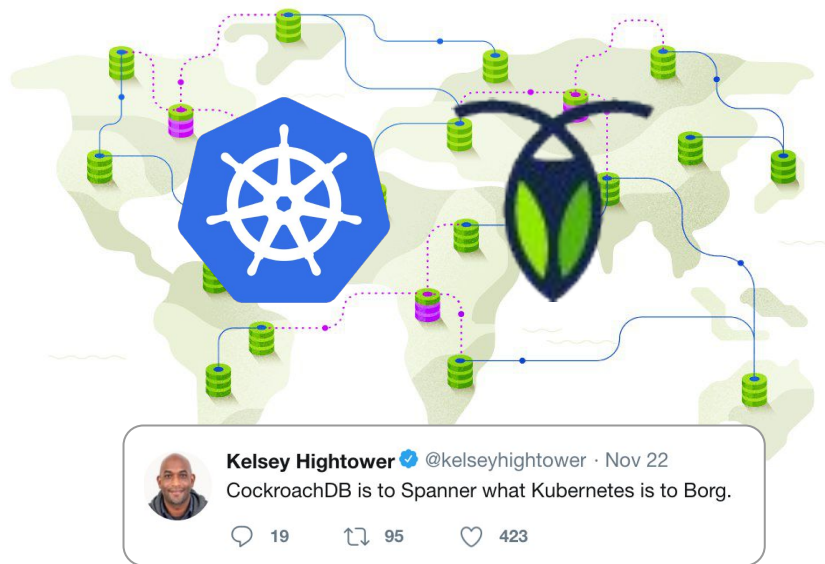
openstack.



CockroachDB and Kubernetes

Common distributed architecture

- Natural fit for pods and orchestration
- Helm chart available
- Multi-region and global scale
- Geotagging within CRDB helps tie compute to data and locality




CockroachDB uses the Storage class and PV claim to mount a volume within a cluster and then builds on stateful sets, so we naturally inherit the controls and power of Kubernetes

Installing CockroachDB on K8s

K8s Installation

- Helm chart or customize the YAML on K8s 1.8+.
- Uses a PV with Stateful Sets.
- Runs best with SSDs.




bit.ly/crdbk8s

Installing CockroachDB on K8s

K8s Installation

bit.ly/crdb-k8ideas

bit.ly/crdb-docker-ideas



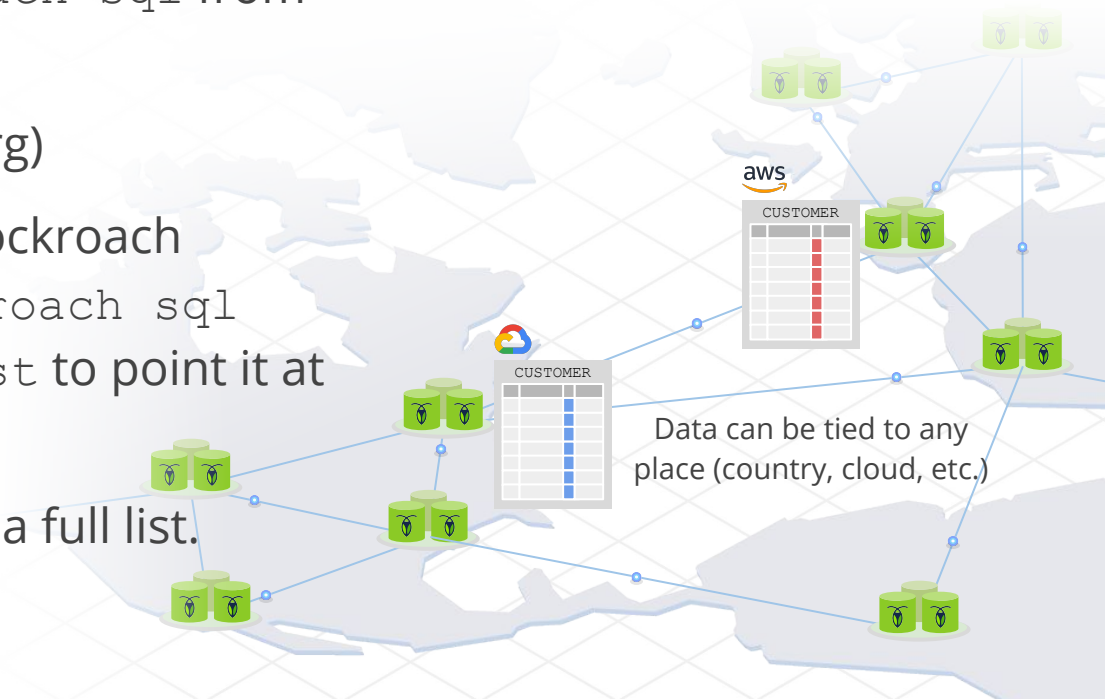
<https://github.com/timveil-cockroach/kubernetes-examples>

<https://github.com/timveil-cockroach/docker-examples>

Connecting to CockroachDB

Connection options for CockroachDB

1. With the SQL shell. (`cockroach sql` from the command line.)
2. With jdbc (jdbc.postgresql.org)
3. Within Kubernetes. (Use a cockroach container and run the `cockroach sql` command and set the `--host` to point it at the internal vip.)
4. See bit.ly/crdb-connect for a full list.





CockroachDB

Cloud native distributed SQL for the cloud native future

Cloud Neutral

Build across on-prem, cloud, hybrid cloud and multi cloud environments



Open Source

UN-opinionated and community driven
so you are not tied to any cloud



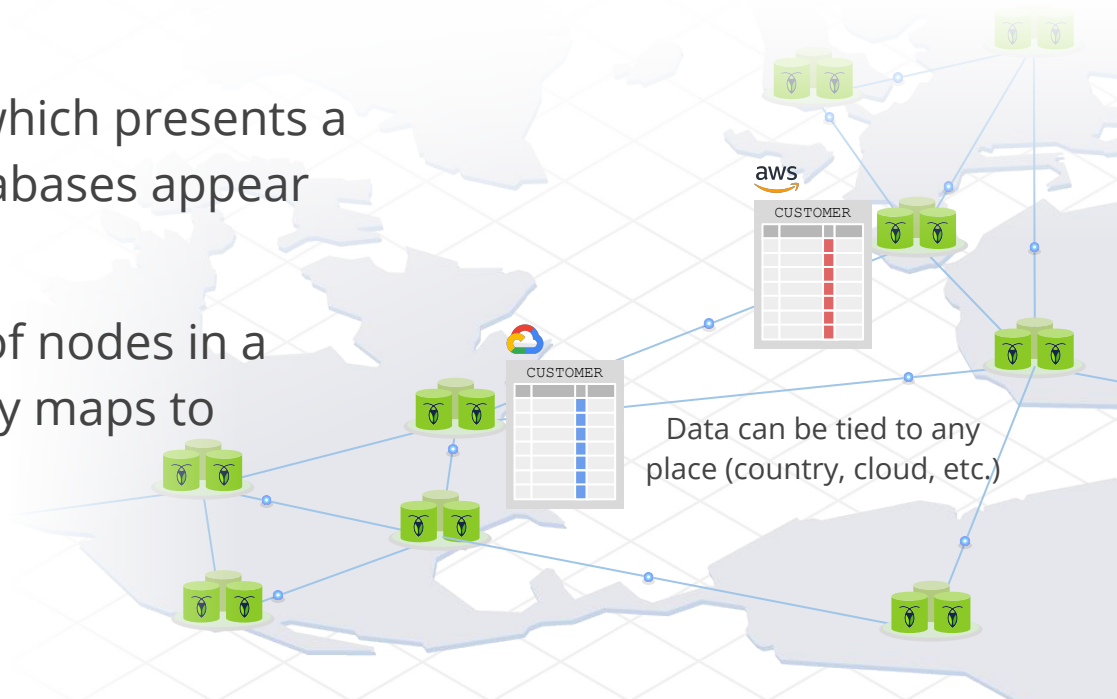
DOWNLOAD NOW!

CockroachDB Architecture Overview



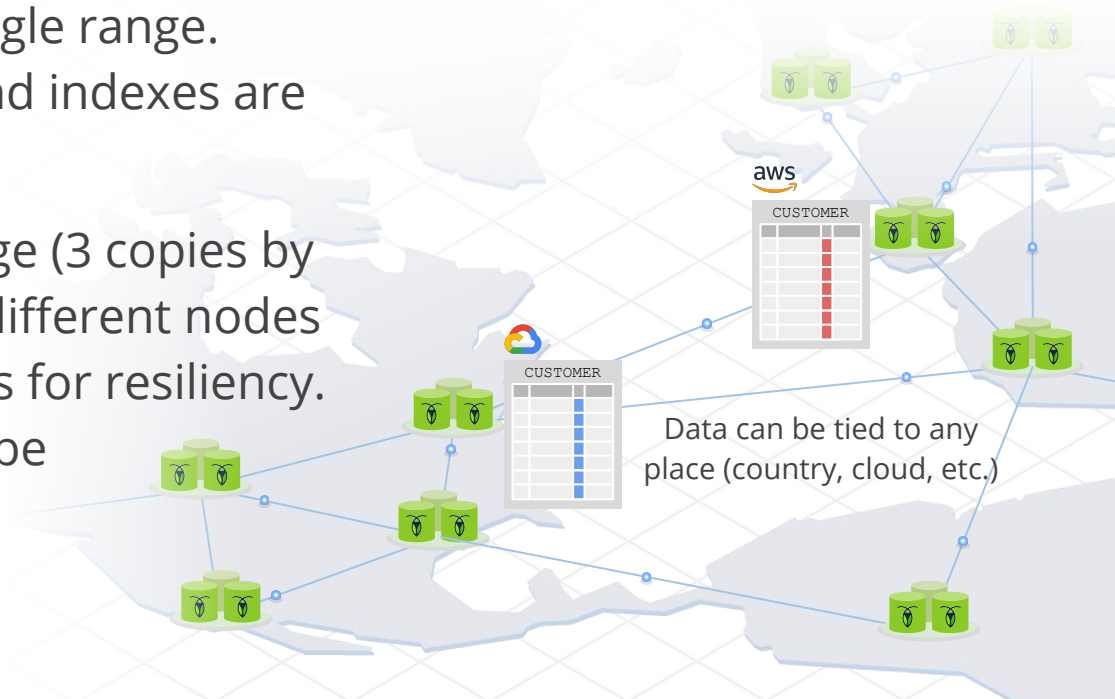
CockroachDB Cluster concepts

1. **Node:** An individual machine running CockroachDB. Many nodes join together to create your cluster.
2. **Cluster:** A group of nodes which presents a single DBMS engine. All databases appear as single logical instances.
3. **Replication Zone:** Groups of nodes in a single location, which usually maps to regions.



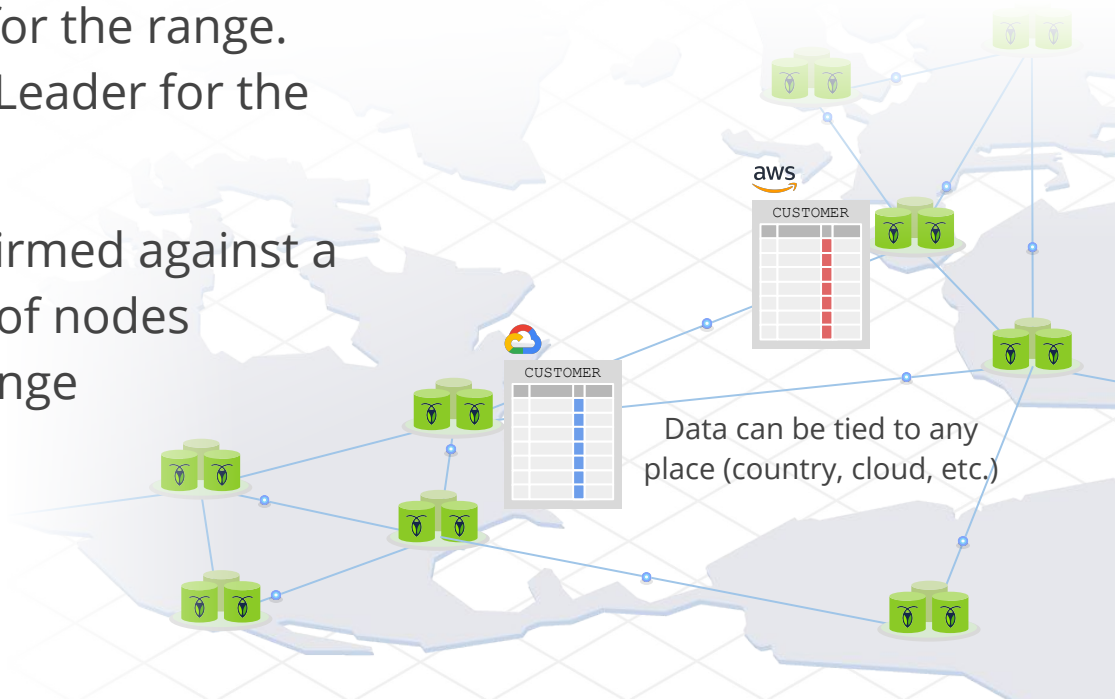
CockroachDB Data concepts

1. **Range:** User Data is stored in ranges, contiguous chunks of data so that every key can always be found in a single range. (64MB by default.) Tables and indexes are spread across ranges.
2. **Replica:** A copy of each range (3 copies by default) which is stored on different nodes in different replication zones for resiliency. The **replication factor** can be configured.

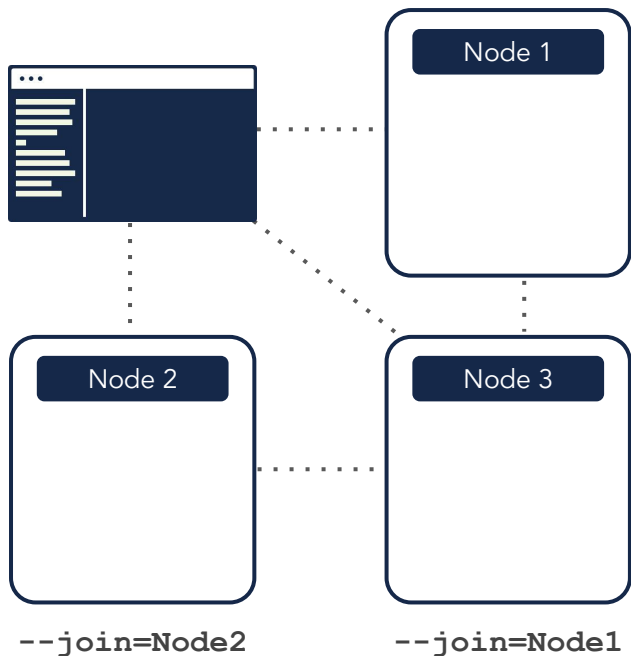


CockroachDB Data concepts

3. **Leaseholder:** Each range designates one one “leaseholder” replica that coordinates all read and write requests for the range. This almost always the Raft Leader for the range.
4. **Consensus:** Writes are confirmed against a consensus when a majority of nodes containing replicas of the range acknowledge the write.



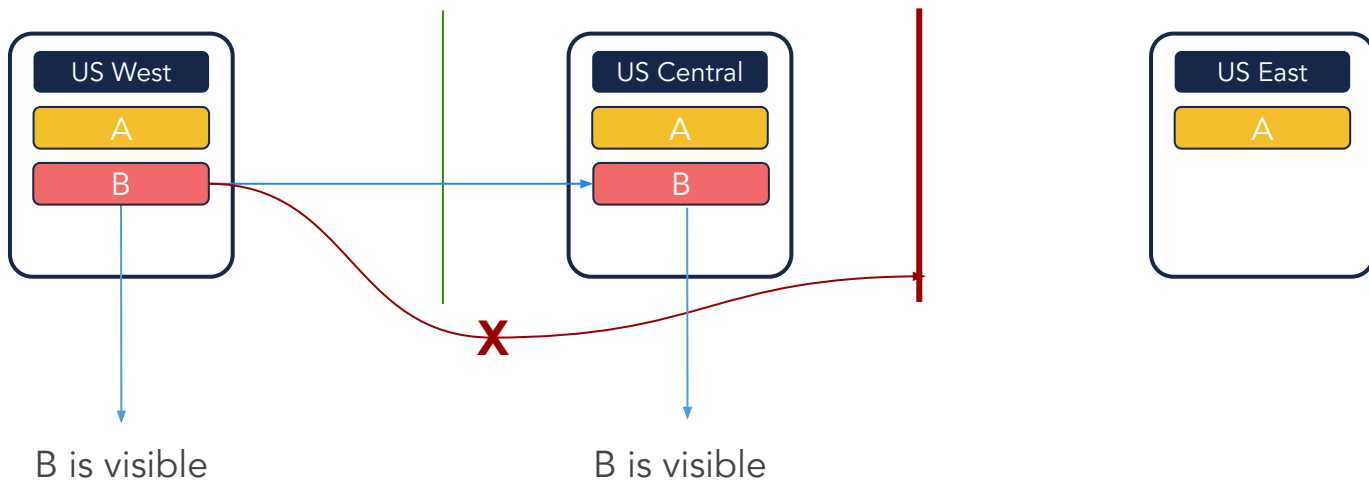
Simple Deployment with Symmetric Nodes



- Download single binary
- Install with a single command
- Every node is a client gateway
- Each node is a share-nothing member of the cluster.
- Applications see one logical DB

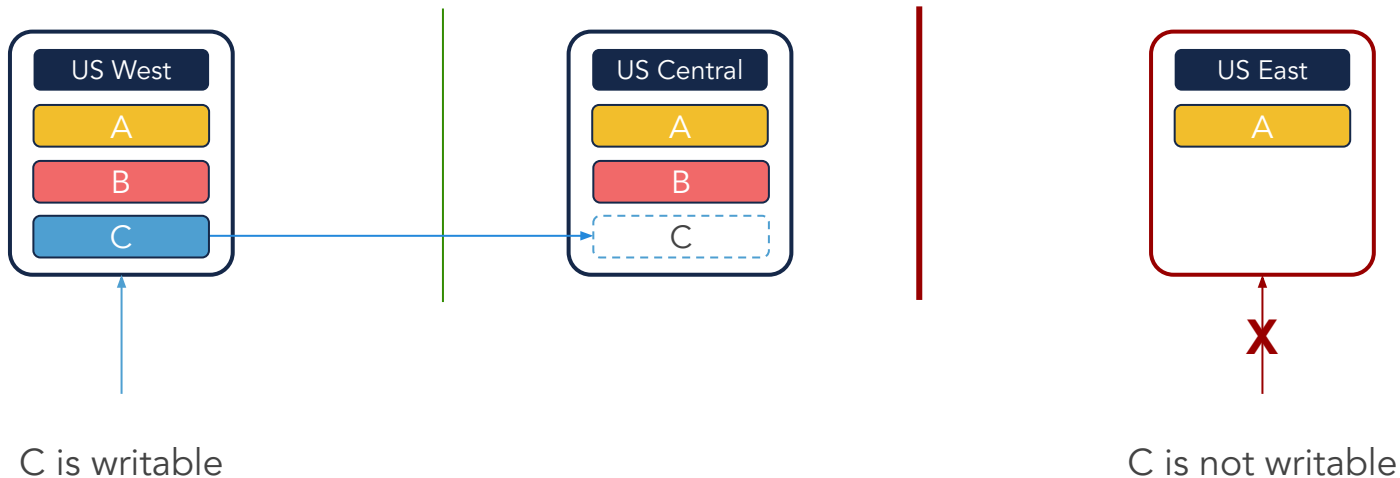
Consistency And Resiliency Between Datacenters

Raft majority consensus means data is always consistent and available.








Consistency Between Datacenters

Majority consensus can also write even if some nodes are down.




Scale Out




- Add additional nodes to scale
- Not limited to one-at-a-time, can add many at once.

Scale Out



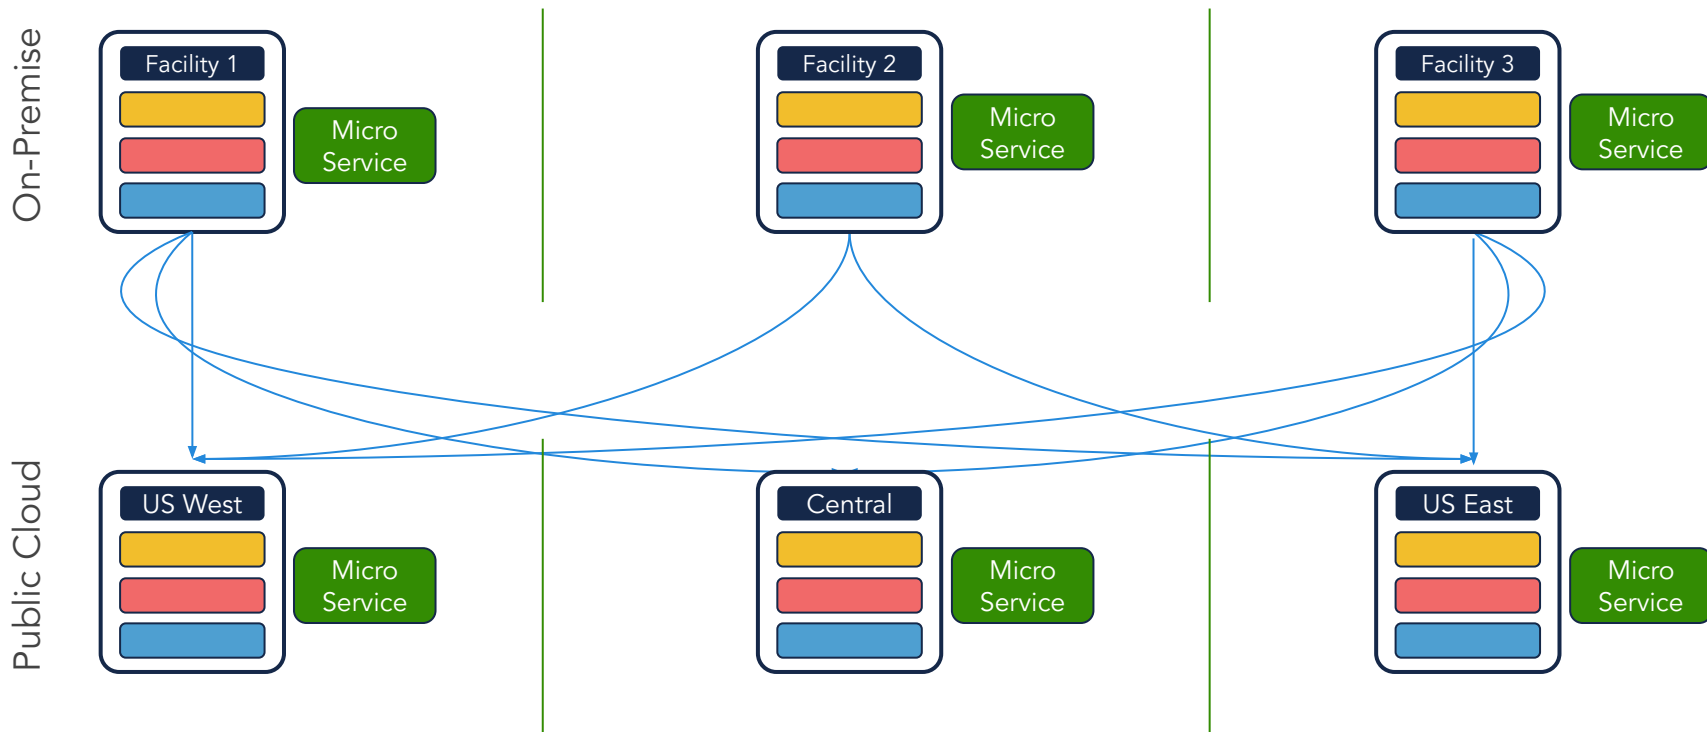
- Nodes self-organize through rebalancing
 - Fill available space from new nodes
 - Automated failover and recovery
 - Invisible to applications

Scale In



- Scale-in by simply shutting down nodes and data will be moved to the rest of the cluster.
- Can remove many nodes at once.

Deploy and Migrate Across Clouds



Underneath the Hood



Photo: [stshank](#)

Multi-version concurrency control (MVCC)

- Values are never updated “in place”, newer versions shadow older versions.
- Tombstones are used to delete values.
- Provides snapshot to each transaction.

Underneath the Hood



Photo: [Taki Steve](#)

Internal Architecture

- Data is stored on disk in a Key-Value database (RocksDB).
- Keys and values are strings lexicographically ordered by key.
- Monolithic key-space.

Underneath the Hood

CockroachDB has access to powerful KV primitives

`Get(key)`

`Put(key, value)`

`ConditionalPut(key, value,
expValue)`

`Scan(startKey, endKey)`

`Del(key)`

KV Store details

- Natural separation for distributed transactions.
- Atomic columns enable dynamic schema change.
- Extendable to additional functionality such as geo-partitioning.
- CRDB retains the efficiency of a KV store but gains a natural ability to distribute data and still remain SQL.

Underneath the Hood

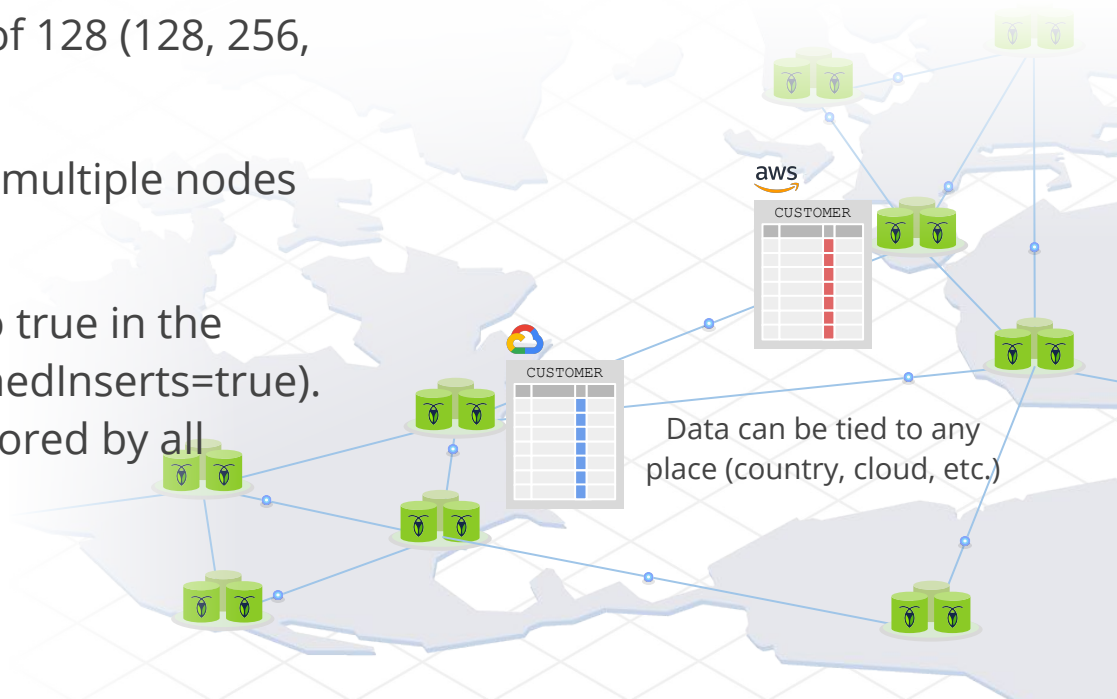
key	value
dog/34/name	carl
dog/34/weight	10.1
dog/7A/name	dagne
dog/7A/weight	13.4
dog/94/name	figment
dog/94/weight	65.8
dog/BC/name	jack
dog/BC/weight	49.7

KV Store details

- All tabular data is stored as monolithic sorted map of KV pairs
- All tables have a primary key
- One key/value pair per column and keys and values are strings
 - Key: <table>/<index>/<key>/<columnName>
 - Value: <columnValue>

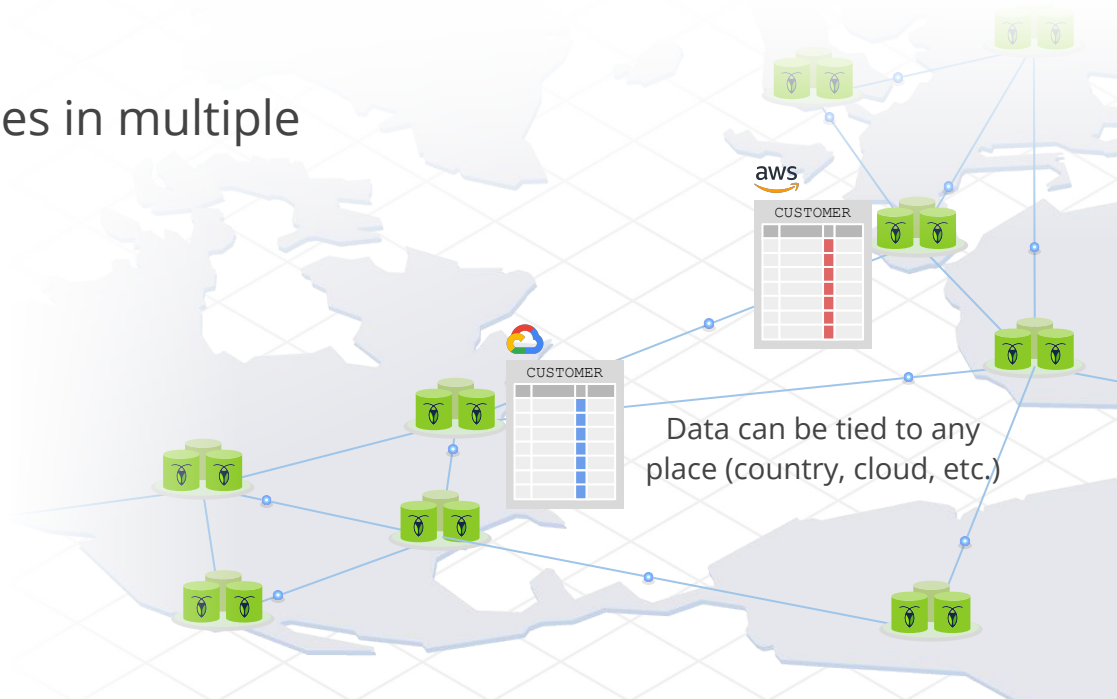
Optimizing performance in Distributed SQL Databases: Bulk Inserts

1. Use the [Import](#) statement for best performance.
2. If you must use Insert, use multi-value inserts with batch sizes in increments of 128 (128, 256, 512, etc).
3. Spread your connection across multiple nodes for higher parallelism.
4. Set the Batched Inserts parm to true in the connection string (reWriteBatchedInserts=true). (But note this is not always honored by all frameworks and drivers.)



Optimizing performance in Distributed SQL Databases: Reads


1. **Leaseholder pinning:** You can pin the leaseholders to specified nodes for faster performance.
2. **Local indexes:** Set up indexes in multiple zones for faster reads.



Connecting to CockroachDB

Connection options for CockroachDB

1. With the SQL shell. (`cockroach sql` from the command line.)
2. With jdbc (jdbc.postgresql.org)
3. Within Kubernetes. (Use a cockroach container and run the `cockroach sql` command and set the `--host` to point it at the internal vip.)
4. See **bit.ly/crdb-connect** for a full list.



Q & A